

Github Link:<https://github.com/Gowtham0504/Phase-2.git>

PROJECT TITLE:*Cracking the Market Code: AI-Driven Stock Price Prediction Using Time Series Analysis*

PHASE-3

1. Problem Statement

Objective:

The goal of this project is to develop a robust AI-driven system capable of predicting short- and mid-term stock price movements using time series analysis techniques. By leveraging deep learning models such as LSTM (Long Short-Term Memory) and Transformer-based architectures, this project aims to identify patterns in historical stock data and generate accurate, data-driven forecasts that can assist traders and investors in making informed decisions.

Problem:

Stock price forecasting remains a highly challenging task due to the volatile, non-linear, and non-stationary nature of financial markets. Traditional statistical models struggle to capture complex temporal dependencies and react to real-time market events. The challenge lies in building a predictive model that can:

- ✓ Effectively learn from historical price and volume data,
- ✓ Incorporate external factors (e.g., news sentiment, macroeconomic indicators),
- ✓ Handle noise and avoid overfitting,

- ✓ Generalize well to unseen market conditions.

Scope:

- ✓ Cracking the Market Code: AI-Driven Stock Price Prediction Using Time Series Analysis
- ✓ Data Source: Historical price data from sources like Yahoo Finance or Alpha Vantage.
- ✓ Target: Predict closing prices or returns for a selected set of stocks.
- ✓ Models: Compare classical time series models (e.g., ARIMA) with AI models (e.g., LSTM, Transformer).
- ✓ Evaluation: Use metrics such as RMSE, MAE, and directional accuracy.

2. Abstract

Introduction:

- Overview of stock market volatility
- Importance of accurate stock price prediction

Traditional Forecasting Methods:

- Brief on statistical models like ARIMA
- Limitations in handling non-linear patterns

Emergence of AI in Finance:

- Role of AI in transforming financial analytics
- Advantages over traditional methods

Understanding Time Series Data:

- Components: trend, seasonality, noise
- Challenges in modeling financial time series

Machine Learning Techniques:

- Supervised vs. unsupervised learning
- Algorithms: Linear Regression, Random Forest, SVM

Deep Learning Approaches:

- Introduction to Neural Networks
- Focus on LSTM networks for sequential data

Data Preprocessing

- Importance of data cleaning
- Handling missing values and outliers

Feature Engineering

- Selecting relevant features
- Incorporating technical indicators

Model Training and Validation

- Splitting data into training and test sets
- Cross-validation techniques

Evaluation Metrics

- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- Mean Absolute Percentage Error (MAPE)

Case Study

- Application of LSTM on historical stock data
- Comparison with traditional models

Challenges and Limitations

- Overfitting in complex models
- Data quality and availability issues

Future Directions

- Incorporating sentiment analysis
- Real-time prediction systems

Conclusion

- Recap of key findings
- Implications for investors and analysts

3. System Requirements

Minimum System Requirements (For Small to Medium-Scale Models)

- **CPU:** Intel Core i5 (10th Gen or later) / AMD Ryzen 5 (3rd Gen or later)
- **RAM:** 16 GB
- **GPU:** NVIDIA GTX 1660 (6GB VRAM) or RTX 2060
- **Storage:** 512 GB SSD
- **OS:** Windows 10 / Ubuntu 20.04 / macOS 12 or later
- **Software/Frameworks:**
 - Python 3.9+
 - Libraries: numpy, pandas, matplotlib, scikit-learn, statsmodels
 - Deep Learning: TensorFlow 2.x or PyTorch
 - Time Series: prophet, tsfresh, sktime
 - IDE: VS Code / JupyterLab
 - Optional: Docker, Anaconda

Recommended System Requirements (For Deep Learning & Faster Training)

- **CPU:** Intel Core i7 (12th Gen or later) / AMD Ryzen 7 (5th Gen or later)
- **RAM:** 32 GB (especially for processing long time series)
- **GPU:** NVIDIA RTX 3060 Ti (8GB VRAM) or better (e.g., RTX 4070/4080 for faster deep learning)
- **Storage:** 1 TB SSD (fast I/O for large datasets)
- **OS:** Ubuntu 22.04 LTS (preferred for ML pipelines) / Windows 11
- **Software/Frameworks:**
 - Python 3.11+
 - Libraries: pandas, numpy, scikit-learn, statsmodels, matplotlib, seaborn
 - Deep Learning: PyTorch (with CUDA 11.x) or TensorFlow 2.15+

- Time Series: prophet, sktime, darts
- Data Handling: dask, modin (if datasets are huge)
- GPU drivers & CUDA toolkit installed
- JupyterLab / VS Code / PyCharm
- Version Control: git

Optional but Useful for Production-Level System

- **Cloud GPU Option:** AWS (EC2 with NVIDIA A100), Google Cloud, Azure ML
- **Database:** PostgreSQL / MySQL / TimescaleDB (for time series data storage)
- **Big Data:** Apache Spark (if dealing with massive stock tick data)
- **Model Deployment:** Docker + FastAPI/Streamlit + Kubernetes (for live prediction service)
- **Monitoring:** MLflow / Weights & Biases

Extra Tools for Stock Market AI Project

- Financial APIs: Yahoo Finance API, Alpha Vantage, IEX Cloud
- Backtesting libraries: backtrader, bt, zipline
- Technical Indicators: ta-lib, finta

4. Objectives

- To collect and preprocess historical stock market data
- Acquire time series data (e.g., open, high, low, close, volume) from financial APIs and clean, normalize, and format it for model training.
- To explore and analyze stock price trends and patterns
- Perform exploratory data analysis (EDA) to understand price movements, detect

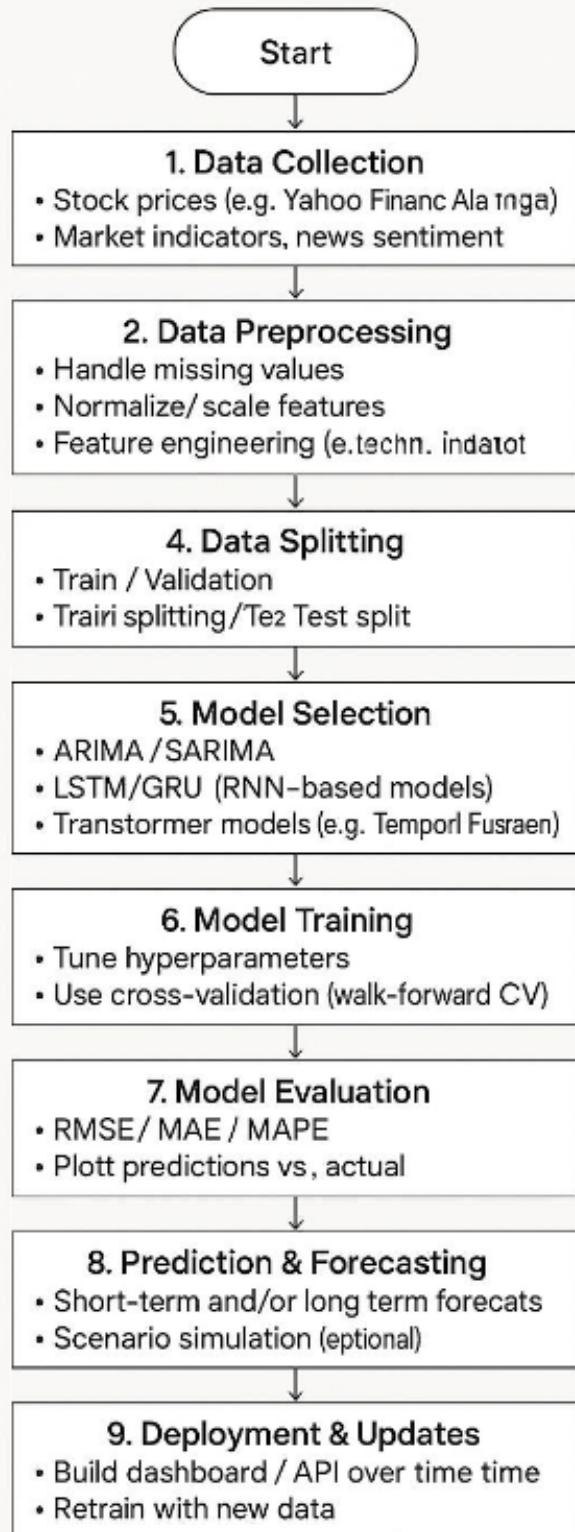
- seasonality, and identify market anomalies.
- To engineer relevant features from time series and external data
- Generate technical indicators (e.g., moving averages, RSI), lagged variables, and
- incorporate sentiment or macroeconomic indicators as additional inputs.
- To implement and compare multiple predictive models Develop and evaluate classical models (ARIMA, SARIMA) and AI models (LSTM, GRU,
- Transformer) for predicting future stock prices.
- To evaluate model performance using appropriate metrics
- Use performance measures like RMSE, MAE, and directional accuracy to assess model effectiveness and ensure reliable forecasts.
- To build a deployable prototype for real-time or batch predictions
- Create a functional pipeline or dashboard that visualizes forecasts and could be adapted for live market data.
- To assess the potential and limitations of AI in financial forecasting

- Analyze model behavior under different market conditions and highlight challenges such as overfitting, data noise, and response to unexpected events.

5. Flowchart of the Project Workflow

AI-Driven Stock Price Prediction

– Time Series Analysis



6. Data Description: AI-Driven Stock Price Prediction (Time Series Analysis)

1. Stock Price Data

Source: Yahoo Finance, Alpha Vantage, Quandl, or similar.

Frequency: Daily (can also use intraday, weekly, or monthly depending on goal).

Fields:

Date: Timestamp of the observation.

Open: Price at market open.

High: Highest price during the day.

Low: Lowest price during the day.

Close: Final price when the market closes.

Adj Close: Adjusted closing price (accounts for splits/dividends).

Volume: Number of shares traded during the day.

2. Derived Features (Optional but recommended)

Technical Indicators (from libraries like TA-Lib or manually engineered):

SMA: Simple Moving Average.

EMA: Exponential Moving Average.

RSI: Relative Strength Index.

MACD: Moving Average Convergence Divergence.

Bollinger Bands: Price volatility bands.

Lagged Values:

Close_{t-1}, Close_{t-2}, etc., to help models capture temporal dependencies.

3. Sentiment or External Data (Optional)

News Headlines or Tweets (with sentiment score using NLP models).

Market Indices (e.g., S&P 500, Nasdaq).

Economic Indicators (e.g., interest rates, inflation data).

4. Target Variable

Usually:

Future Close Price (regression)

Price Movement Direction (classification: up/down)

7. Data Processing: AI-Driven Stock Price Prediction (Time Series Analysis)

1. Data Cleaning

- Remove duplicates: Ensure each date has only one row.

- Handle missing values:
- Fill with forward-fill (ffill), backward-fill (bfill), or interpolation.
- Drop rows if missing critical data like prices or volume.

2. Date Parsing & Indexing

- Convert Date column to datetime format.
- Set Date as the index for time series modeling.

3. Feature Engineering

- Lag features: Include past values as features (e.g., Close_t-1, Close_t-2, ...).
- Rolling statistics:
 - Moving average (e.g., 7-day, 14-day, 30-day)
 - Rolling standard deviation or volatility
- Technical indicators:
 - SMA, EMA, RSI, MACD, Bollinger Bands
 - Use ta-lib or pandas-ta for easy computation

4. Target Variable Creation

- Regression: Predict future Close price (e.g., Close_t+1, Close_t+5)
- Classification: Predict direction (1 for up, 0 for down)

5. Data Normalization/Scaling

- Use MinMaxScaler or StandardScaler (especially important for neural networks like LSTM)

- Apply scaling to features only, not the date or target variable during model training.

6. Train-Test Split (Chronologically)

- Use 70–80% for training, remaining for testing or validation.
- Avoid random splitting to preserve temporal integrity.



1. Load the Data

Make sure your data includes:

- Date/Time (Index)
- Open, High, Low, Close, Volume
- Possibly technical indicators (optional at this stage)



2. Basic Statistical Summary

Understand the central tendencies and spread of the price and volume data. A small icon depicting a bar chart, representing data visualization.

3. Time Series Plots

Visualize the trend and seasonality.



4. Rolling Statistics (Moving Averages)

Helps smooth out fluctuations and identify trend direction.

5. Volatility Analysis

Volatility can signal risk and opportunity.

6. Correlation Heatmap

Evaluate relationships between numerical columns.

7. Stationarity Check (ADF Test)

Stationary data is key for many models (ARIMA, SARIMA).

Next Steps After EDA:

- Feature Engineering (e.g., RSI, MACD)
- Train-test split
- Normalize/scale data
- Apply models: LSTM, Prophet, ARIMA, XGBoost

8. Feature Engineering for Stock Price Prediction

Steps for EDA in Stock Price Time Series Analysis

1. Load the Data

- Use stock price datasets (from Yahoo Finance, Alpha Vantage, etc.).
- Common columns: Date, Open, High, Low, Close, Adj Close, Volume.

```
import pandas as pd  
df = pd.read_csv('stock_prices.csv',  
parse_dates=['Date'], index_col='Date')  
print(df.head())
```

2. Check Missing Values & Data Types

- Stock datasets often have missing data on weekends/holidays.

```
print(df.info())  
print(df.isnull().sum())
```

3. Visualize Stock Price Trends (Line Plots)

- Plot the **closing price** over time to see overall trends.

```
import matplotlib.pyplot as plt  
df['Close'].plot(figsize=(12,6), title='Stock Closing  
Price Over Time')  
plt.show()
```

4. Moving Averages (Smoothing)

- Use *Simple Moving Averages* (SMA) to smooth short-term fluctuations.

`df['SMA50'] = df['Close'].rolling(window=50).mean()`

`df['SMA200'] =`
`df['Close'].rolling(window=200).mean()`

`df[['Close', 'SMA50', 'SMA200']].plot(figsize=(12,6))`

`plt.show()`

5. Volatility Analysis

- Use daily returns or percentage changes to analyze stock volatility.

`df['Daily Return'] = df['Close'].pct_change()`

`df['Daily Return'].hist(bins=50, figsize=(10,6))`

`plt.show()`

6. Seasonality & Trend Decomposition

- Decompose time series into **trend**, **seasonality**, and **residual**.

`from statsmodels.tsa.seasonal import
seasonal_decompose`

`result = seasonal_decompose(df['Close'],
model='multiplicative', period=30)`

`result.plot()`

`plt.show()`

7. Correlation Heatmap

- Correlation between *Open*, *High*, *Low*, *Close*, *Volume*.

`import seaborn as sns`

```
plt.figure(figsize=(8,6))
```

```
sns.heatmap(df.corr(), annot=True,  
cmap='coolwarm')
```

```
plt.show()
```

8. Lag Plots (Autocorrelation)

- To check if past values influence future ones (good for time series models).

```
from pandas.plotting import lag_plot
```

```
lag_plot(df['Close'])
```

```
plt.show()
```

9. Stationarity Test (Augmented Dickey-Fuller Test)

- Important before applying ARIMA/LSTM models.

```
from statsmodels.tsa.stattools import adfuller
```

```
result = adfuller(df['Close'].dropna())
print(f'ADF Statistic: {result[0]}')
print(f'p-value: {result[1]}')
```

10. Volume vs Price Analysis

- Check if high trading volumes impact stock price changes.

```
df[['Volume', 'Close']].plot(subplots=True,
figsize=(10,6))
plt.show()
```

9. Feature Engineering for Stock Price Prediction

1. Lag Features

These are past values used to predict the future.

2. Rolling Window Statistics

These show trends and volatility over time.



3. Returns

Helps identify momentum or reversals.

4. Technical Indicators

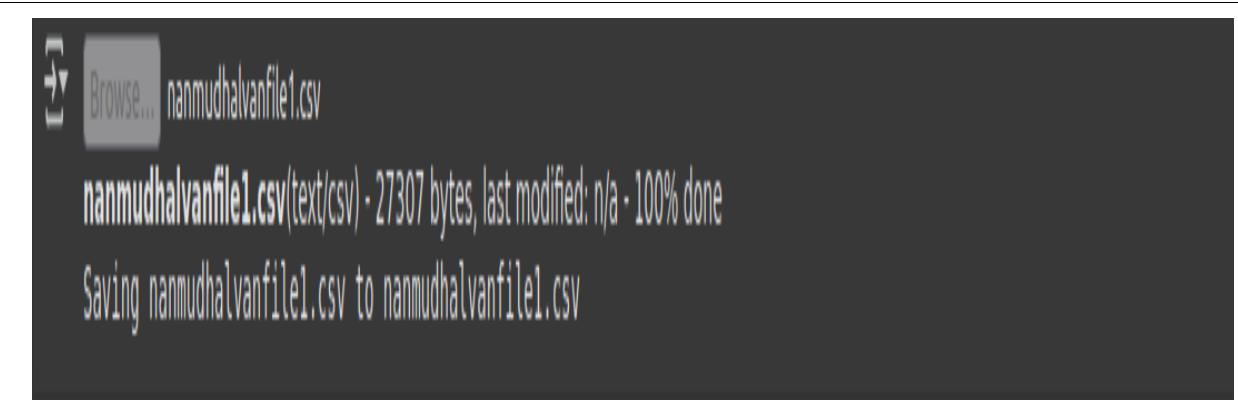
- Popular indicators used in trading.
- Relative Strength Index (RSI)
- Moving Average Convergence Divergence (MACD)
- Bollinger Bands

3. Target Variable for Prediction

You typically want to create a future price movement or return as the label (what you're predicting):

4. Prepare for ML Model

- Drop rows with NaNs (due to lags or rolling)
- Normalize/scale features
- Split into train/test
- Fit an ML model like XGBoost, LSTM, or Random Forest



	Date	Symbol	Series	Prev Close	Open	High	Low	Last	Close	VWAP	Volume	Turnover	Trades	Deliverable Volume	%Deliverble
0	01/01/15	INFY	EQ	1972.55	1968.95	1982.00	1956.9	1971.00	1974.40	1971.34	500691	987000000000000	14908	258080	0.5154
1	01/02/15	INFY	EQ	1974.40	1972.00	2019.05	1972.0	2017.95	2013.20	2003.25	1694580	339000000000000	54166	1249104	0.7371
2	01/05/15	INFY	EQ	2013.20	2009.90	2030.00	1977.5	1996.00	1995.90	2004.59	2484256	498000000000000	82694	1830962	0.7370
3	01/06/15	INFY	EQ	1995.90	1980.00	1985.00	1934.1	1965.10	1954.20	1954.82	2416829	472000000000000	108209	1772070	0.7332
4	01/07/15	INFY	EQ	1954.20	1965.00	1974.75	1950.0	1966.05	1963.55	1962.59	1812479	356000000000000	62463	1317720	0.7270

```
Epoch 1/10
6/6 1s 50ms/step - loss: 0.1909
Epoch 2/10
6/6 1s 50ms/step - loss: 0.0360
Epoch 3/10
6/6 1s 49ms/step - loss: 0.0283
Epoch 4/10
6/6 1s 49ms/step - loss: 0.0242
Epoch 5/10
6/6 0s 53ms/step - loss: 0.0283
Epoch 6/10
6/6 0s 50ms/step - loss: 0.0167
Epoch 7/10
6/6 1s 53ms/step - loss: 0.0189
Epoch 8/10
6/6 1s 50ms/step - loss: 0.0159
Epoch 9/10
6/6 0s 55ms/step - loss: 0.0098
Epoch 10/10
6/6 0s 49ms/step - loss: 0.0121
<keras.callbacks.history.History at 0x78b40b78650>
```

10. Model Building

Step 1: Choose the Type of Model

For time series stock prediction, popular models include:

Step 2: Data Preprocessing

- Use the *Close* price for prediction.
- Normalize/scale the data.
- Split into sequences (time steps).

import numpy as np

from sklearn.preprocessing import MinMaxScaler

data = df[['Close']].values

scaler = MinMaxScaler()

scaled_data = scaler.fit_transform(data)

def create_sequences(data, time_steps=60):

X, y = [], []

for i in range(time_steps, len(data)):

X.append(data[i-time_steps:i, 0])

y.append(data[i, 0])

return np.array(X), np.array(y)

`X, y = create_sequences(scaled_data)`

`X = X.reshape((X.shape[0], X.shape[1], 1))`

 **Step 3: Train-Test Split**

`split = int(0.8 * len(X))`

`X_train, X_test = X[:split], X[split:]`

`y_train, y_test = y[:split], y[split:]`

 **Step 4: Build the LSTM Model**

`import tensorflow as tf`

`from tensorflow.keras.models import Sequential`

`from tensorflow.keras.layers import LSTM, Dense, Dropout`

`model = Sequential([`

`LSTM(50, return_sequences=True, input_shape=(X_train.shape[1], 1)),`

`Dropout(0.2),`

`LSTM(50, return_sequences=False),`

`Dropout(0.2), Dense(1)])`

`model.compile(optimizer='adam', loss='mean_squared_error')`

`model.summary()`

Step 5: Train the Model

```
history = model.fit(X_train, y_train, epochs=20, batch_size=32,  
validation_data=(X_test, y_test))
```

Step 6: Make Predictions and Invert Scaling

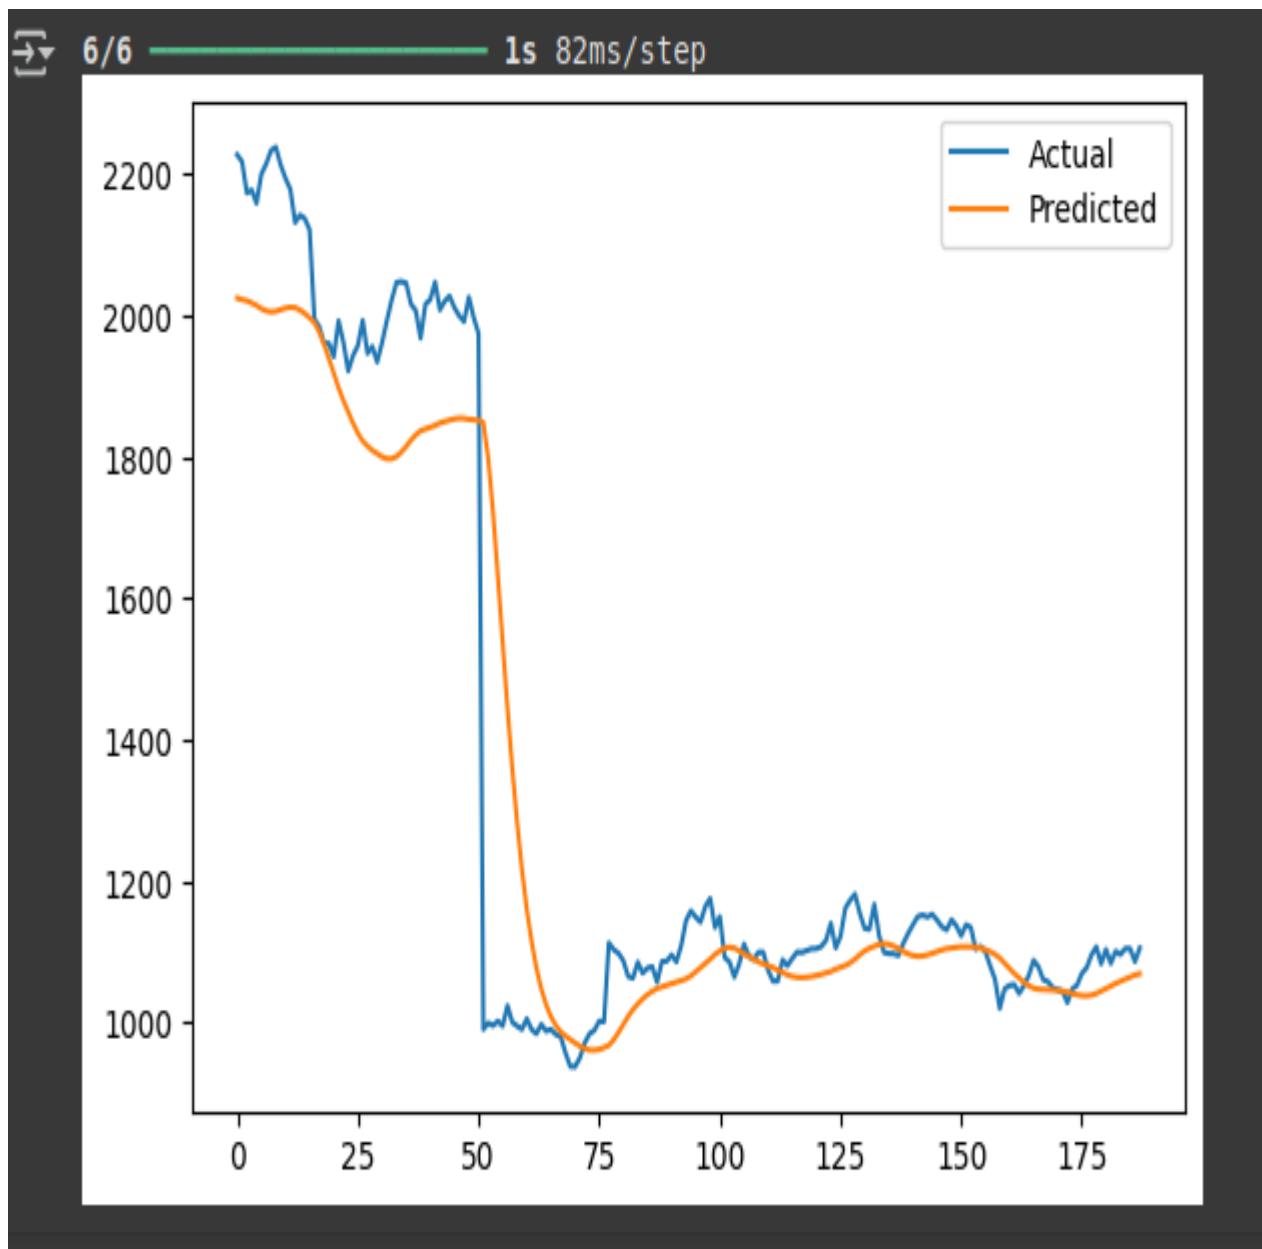
```
predicted = model.predict(X_test)  
predicted_prices = scaler.inverse_transform(predicted.reshape(-1,  
1))  
real_prices = scaler.inverse_transform(y_test.reshape(-1, 1))
```

Step 7: Plot the Predictions

```
import matplotlib.pyplot as plt  
plt.figure(figsize=(12,6))  
plt.plot(real_prices, color='blue', label='Actual Stock Price')  
plt.plot(predicted_prices, color='red', label='Predicted Stock Price')  
plt.title('Stock Price Prediction')  
plt.xlabel('Time')  
plt.ylabel('Price')  
plt.legend()
```

plt.show()

11. Model Evaluation



```
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas<3.0,>=1.0->gradio) (2025.2)
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (0.7.0)
Requirement already satisfied: pydantic-core==2.33.1 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (2.33.1)
Requirement already satisfied: typing-inspection>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from pydantic<2.12,>=2.0->gradio) (0.4.0)
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (8.1.8)
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (1.5.4)
Requirement already satisfied: rich>=10.11.0 in /usr/local/lib/python3.11/dist-packages (from typer<1.0,>=0.12->gradio) (13.9.4)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas<3.0,>=1.0->gradio) (1.17.0)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich>=10.11.0->typer<1.0,>=0.12->gradio) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich>=10.11.0->typer<1.0,>=0.12->gradio) (2.19.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface-hub>=0.28.1->gradio) (3.4.1)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->huggingface-hub>=0.28.1->gradio) (2.4.0)
Requirement already satisfied: moudl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich>=10.11.0->typer<1.0,>=0.12->gradio) (0.1.2)
Downloading gradio-5.29.0-py3-none-any.whl (54.1 MB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 54.1/54.1 MB 8.1 MB/s eta 0:00:00
Downloading gradio_client-1.10.0-py3-none-any.whl (322 kB)
  ━━━━━━━━━━━━━━━━ 322.9/322.9 kB 25.2 MB/s eta 0:00:00
Downloading aiofiles-24.1.0-py3-none-any.whl (15 kB)
Downloading fastapi-0.115.12-py3-none-any.whl (95 kB)
  ━━━━━━━━━━ 95.2/95.2 kB 8.9 MB/s eta 0:00:00
Downloading groovy-0.1.2-py3-none-any.whl (14 kB)
Downloading python_multipart-0.0.20-py3-none-any.whl (24 kB)
Downloading ruff-0.11.8-py3-none-manylinux_2_17_x86_64_manylinux2014_x86_64.whl (11.5 MB)
  ━━━━━━━━━━ 11.5/11.5 MB 115.8 MB/s eta 0:00:00
Downloading safehttpx-0.1.6-py3-none-any.whl (8.7 kB)
Downloading semantic_version-2.10.0-py2.py3-none-any.whl (15 kB)
Downloading starlette-0.46.2-py3-none-any.whl (72 kB)
  ━━━━━━━━ 72.0/72.0 kB 6.9 MB/s eta 0:00:00
Downloading tomkit-0.13.2-py3-none-any.whl (37 kB)
Downloading uvicorn-0.34.2-py3-none-any.whl (62 kB)
  ━━━━━━━━ 62.5/62.5 kB 4.8 MB/s eta 0:00:00
Downloading ffmpy-0.5.0-py3-none-any.whl (6.0 kB)
Downloading pydub-0.25.1-py2.py3-none-any.whl (32 kB)
Installing collected packages: pydub, uvicorn, tomkit, semantic-version, ruff, python-multipart, groovy, ffmpy, aiofiles, starlette, safehttpx, gradio-client, fastapi, gradio
Successfully installed aiofiles-24.1.0 fastapi-0.115.12 ffmpy-0.5.0 gradio-5.29.0 gradio-client-1.10.0 groovy-0.1.2 pydub-0.25.1 python-multipart-0.0.20 ruff-0.11.8 safehttpx-0.1.6 semantic-version-2.10.0 st
It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sharing must be enabled. Automatically setting 'share=True' (you can turn this off by setting 'share=False' in

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()
* Running on public URL: https://fbc9bbb18ae37ebb78.gradio.live
```

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run 'gradio deploy' from the terminal in the working directory to deploy to Hugging Face Spaces (<https://huggingface.co/spaces>)

```
[+] Collecting gradio
  Downloading gradio-5.29.0-py3-none-any.whl.metadata (16 kB)
Collecting aiofiles<25.0,>=22.0 (from gradio)
  Downloading aiofiles-24.1.0-py3-none-any.whl.metadata (10 kB)
Requirement already satisfied: anyio<5.0,>=3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.9.0)
Collecting fastapi<1.0,>=0.115.2 (from gradio)
  Downloading fastapi-0.115.12-py3-none-any.whl.metadata (27 kB)
Collecting ffdpx (from gradio)
  Downloading ffdpx-0.5.0-py3-none-any.whl.metadata (3.0 kB)
^Collecting gradio-client==1.10.0 (from gradio)
  Studio Downloading gradio_client-1.10.0-py3-none-any.whl.metadata (7.1 kB)
Collecting groovy=>0.1 (from gradio)
  Downloading groovy-0.1.2-py3-none-any.whl.metadata (6.1 kB)
Requirement already satisfied: httpx>=0.24.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.28.1)
Requirement already satisfied: huggingface-hub>=0.28.1 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.38.2)
Requirement already satisfied: jinja2<4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.1.6)
Requirement already satisfied: markupsafe<4.0,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.0.2)
Requirement already satisfied: numpy<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.0.2)
Requirement already satisfied: orjson=>3.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (3.10.17)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from gradio) (24.2)
Requirement already satisfied: pandas<3.0,>=1.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.2.2)
Requirement already satisfied: pillow<2.0,>=8.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (11.2.1)
Requirement already satisfied: pydantic<2.12,>=2.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (2.11.3)
Collecting pydub (from gradio)
  Downloading pydub-0.25.1-py2.py3-none-any.whl.metadata (1.4 kB)
Collecting python-multipart>=0.0.18 (from gradio)
  Downloading python_multipart-0.0.20-py3-none-any.whl.metadata (1.8 kB)
Requirement already satisfied: pyyaml<7.0,>=5.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (6.0.2)
Collecting ruff>=0.9.3 (from gradio)
  Downloading ruff-0.11.8-py3-none-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (25 kB)
Collecting safetensors<0.2.0,>=0.1.6 (from gradio)
  Downloading safetensors-0.1.6-py3-none-any.whl.metadata (4.2 kB)
Collecting semantic-version=>2.0 (from gradio)
  Downloading semantic_version-2.10.0-py2.py3-none-any.whl.metadata (9.7 kB)
Collecting starlette<1.0,>=0.40.0 (from gradio)
  Downloading starlette-0.46.2-py3-none-any.whl.metadata (6.2 kB)
Collecting tomkit<0.14.0,>=0.12.0 (from gradio)
  Downloading tomkit-0.13.2-py3-none-any.whl.metadata (2.7 kB)
Requirement already satisfied: typer<1.0,>=0.12 in /usr/local/lib/python3.11/dist-packages (from gradio) (0.15.3)
Requirement already satisfied: typing-extensions=>4.0 in /usr/local/lib/python3.11/dist-packages (from gradio) (4.13.2)
Collecting uvicorn=>0.14.0 (from gradio)
  Downloading uvicorn-0.34.2-py3-none-any.whl.metadata (6.5 kB)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.0->gradio) (2023.3.2)
Requirement already satisfied: websockets<16.0,>=10.0 in /usr/local/lib/python3.11/dist-packages (from gradio-client==1.10.0->gradio) (15.0.1)
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.11/dist-packages (from anyio<5.0,>=3.0->gradio) (3.10)
```

12. Deployment

- **Deployment Method:** Gradio Interface
- **Public Link:** <https://b3b6b9e55e8e768d27.gradio.live>

UI Screenshot:

The screenshot shows a dark-themed user interface for an AI Stock Price Predictor. At the top center, it says "AI Stock Price Predictor". Below that, a instruction text reads "Enter the last 60 days of closing prices to predict the next day's price.". On the left, there is a section labeled "prices" containing a text input field with "Close Prices" and a "Clear" button. To the right of this is a "Submit" button. On the far right, there is a section labeled "output" with a text input field showing "0" and a "Flag" button below it.

13. Source Code

1. Upload the Dataset

```
from google.colab import files
uploaded = files.upload()
```

2. Load the Dataset

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
df = pd.read_csv('nanmudhalvanfile1.csv')
```

3. Data Exploration

```
df.head()
```

4. Visualize a Few Features

```
import pandas as pd
import numpy as np
df = pd.read_csv('nanmudhalvanfile1.csv')
df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
```

```

scaled_data = scaler.fit_transform(df[['Close']])

def create_sequences(data, seq_length):
    X, y = [], []
    for i in range(seq_length, len(data)):
        X.append(data[i-seq_length:i, 0])
        y.append(data[i, 0])
    return np.array(X), np.array(y)
sequence_length = 60
X, y = create_sequences(scaled_data, sequence_length)
X = np.reshape(X, (X.shape[0], X.shape[1], 1))

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
model = Sequential()
model.add(LSTM(50, return_sequences=True, input_shape=(X.shape[1], 1)))
model.add(LSTM(50))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mean_squared_error')

model.fit(X, y, epochs=10, batch_size=32)

predicted = model.predict(X)
predicted_prices = scaler.inverse_transform(predicted.reshape(-1, 1))
import matplotlib.pyplot as plt
actual = scaler.inverse_transform(y.reshape(-1, 1))
plt.plot(actual, label='Actual')
plt.plot(predicted_prices, label='Predicted')
plt.legend()
plt.show()

```

5. Deployment-Building an Interactive App

```

pip install gradio
import gradio as gr
def predict_next(prices):
    input_seq = scaler.transform(np.array(prices).reshape(-1, 1))

```

```

X_input = np.reshape(input_seq[-60:], (1, 60, 1))
pred = model.predict(X_input)
return scaler.inverse_transform(pred)[0][0]
interface = gr.Interface(
    fn=predict,
    inputs=gr.Dataframe(headers=["Close Prices"], row_count=60),
    outputs="number",
    title="AI Stock Price Predictor",
    description="Enter the last 60 days of closing prices to predict the next day's
                price."
)
interface.launch()

```

13. Team Members and Roles

GOKUL – PROBLEM STATEMENT AND PROJECT

OBJECTIVES

GOWTHAM – FLOWCHART OF THE PROJECT

WORKFLOW, DATA DESCRIPTION, DATA PREPROSSING,

EXPLORATORY DATA ANALYSIS

GOWRI SANKAR – FEATURE ENGGINEERING AND

MODEL BUILDING

HARI PRRASAD – VISUALIZATION OF RESULTS &

MODEL INSIGHTS AND TOOLS & TECHNOLOGIES USED